Autodecision

Should attendance go optional now that AI tutoring beats lecture?

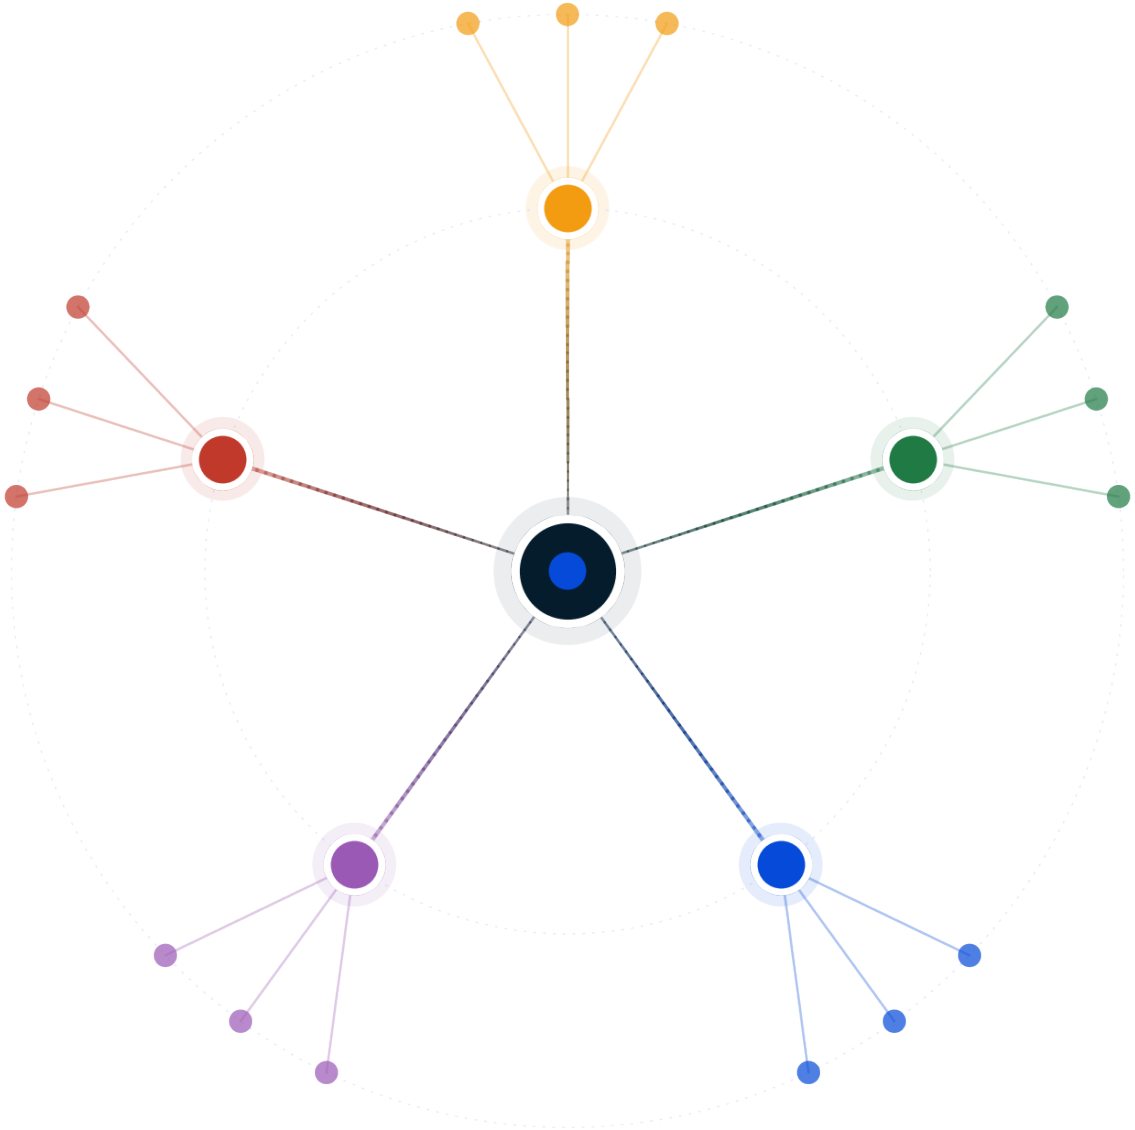

Decision Brief — possibility map and recommendation

April 2026

GENERATED BY AUTODECISION — possibility map with synthesized recommendation, not a single-shot answer

Table of contents

Section	Slides
Recommendation	3
How we explored this decision	4 – 7
What the council found	8 – 13
Path forward	14 – 15
Appendix — Sources	16

Universal optional attendance is strictly dominated; ship a 3-year department pilot with mandatory first-year and dual-track compliance

Do NOT make attendance universally optional. Run a 3-year opt-in pilot in 1-2 departments + mandatory first-year residential + dual-track F-1/Title IV compliance with required in-person seminars + faculty senate and union as co-design partners + Harvard RCT-style AI-hybrid teaching model.

CONFIDENCE

HIGH that universal optional is dominated · MEDIUM-HIGH that the phased-structured path is executable · MEDIUM on regulatory structuring viability

DEPENDS ON

Required in-person seminars satisfy F-1 SEVP and Title IV RSI · Faculty union accepts co-design · Pilot data credible after 3 years · Mandatory first-year preserves first-gen retention · Harvard RCT replicates at institution scale

MONITOR

DOE guidance on RSI · SEVP guidance on in-person seminar components · Faculty union calendar · Pilot cohort year-1 outcomes · First-gen retention vs control · Donor sentiment quarterly · ChatGPT-6 capability benchmarks

PRE-MORTEM

Faculty strikes mid-semester (P=0.15) · Pilot inconclusive from selection bias (P=0.35) · F-1 structured compliance fails legal review (P=0.30) · First-gen dropout despite first-year (P=0.20) · Mid-tier positioning fails (P=0.30)

DEEPEST DISAGREEMENT

Optimist: 60% Harvard RCT replicates at institution scale [C1:optimist]. Pessimist: 25% [C1:pessimist]. Range 25-60% on whether AI-hybrid pedagogy holds across 30K students. Disagreement IS the data.

REVIEW TRIGGER

DOE issues negative RSI guidance · SEVP clarifies one-online-course rule strictly · Peer institution announces similar pilot · Pilot year-1 shows negative equity · Faculty union files grievance · ChatGPT-6 obsoletes current AI-hybrid design

Source: Decision Brief (full mode, 2 iterations, 5/5 council endorsement on h6) · [G1], [C3:regulator], [C4:regulator], [C13:iter2council]

PART 2

How we explored this decision

A 5-persona council, 6 competing hypotheses, two iterations to convergence, 12 adversarial stress scenarios — what survived?

The analysis explored 6 hypotheses across 2 iterations of a 5-persona council, stress-tested by 12 adversarial scenarios

- 5 personas spawned as separate subagents — Optimist, Pessimist, Competitor, Regulator, Customer — each with independent context
- 6 hypotheses generated, ranging from universal optional (h1) to phased-structured pilot (h6 — promoted in iteration 2 as the synthesis winner)
- 2 iterations — convergence reached at iteration 2 with 84% [C1:council] assumption stability and 1 productive contradiction remaining
- Adversary phase generated 12 stress scenarios across 4 worst-cases, 4 black swans, and 4 irrational actors
- Sensitivity analysis on 8 key assumptions — 4 high-sensitivity, 4 medium
- Real-world anchor: Harvard 2025 RCT [G1] showed AI-tutoring effect size 0.73-1.3 SD; the recommendation operationalizes a Harvard-style AI-hybrid teaching model inside a structured pilot.

Source: Decision Brief Convergence Log; engine-protocol.md

Six hypotheses surfaced; only h6 (phased + structured) survived 5/5 council endorsement and adversarial stress

#	Hypothesis	Status	Key assumptions
1	Optional attendance improves outcomes	Conditional	AI tutoring outperforms · Students self-direct · Faculty adapts
2	Attendance collapse degrades social product	Real for elite tier	Peer network IS real product · Tuition premium for network
3	Faculty revolt / union path	Risk if unilateral	Union contracts strict · Faculty identity load-bearing
4	Enrollment bifurcation (mid-tier vulnerable)	Real structure	Elite resists · Middle squeezed · Online MBA precedent
5	De facto already optional	Lecture-only	Revealed preference matches policy · 50-70% [G6] skip routinely
6	NEW: Phased + structured pilot with safeguards	Leading · 5/5 council	Pilot bounds risk · Regulatory structuring viable · Faculty partnership

Source: Decision Brief — Hypotheses Explored

Plotted on reversibility × regulatory exposure, h6 sits in the manageable quadrant; h1 sits where every adversarial scenario amplifies

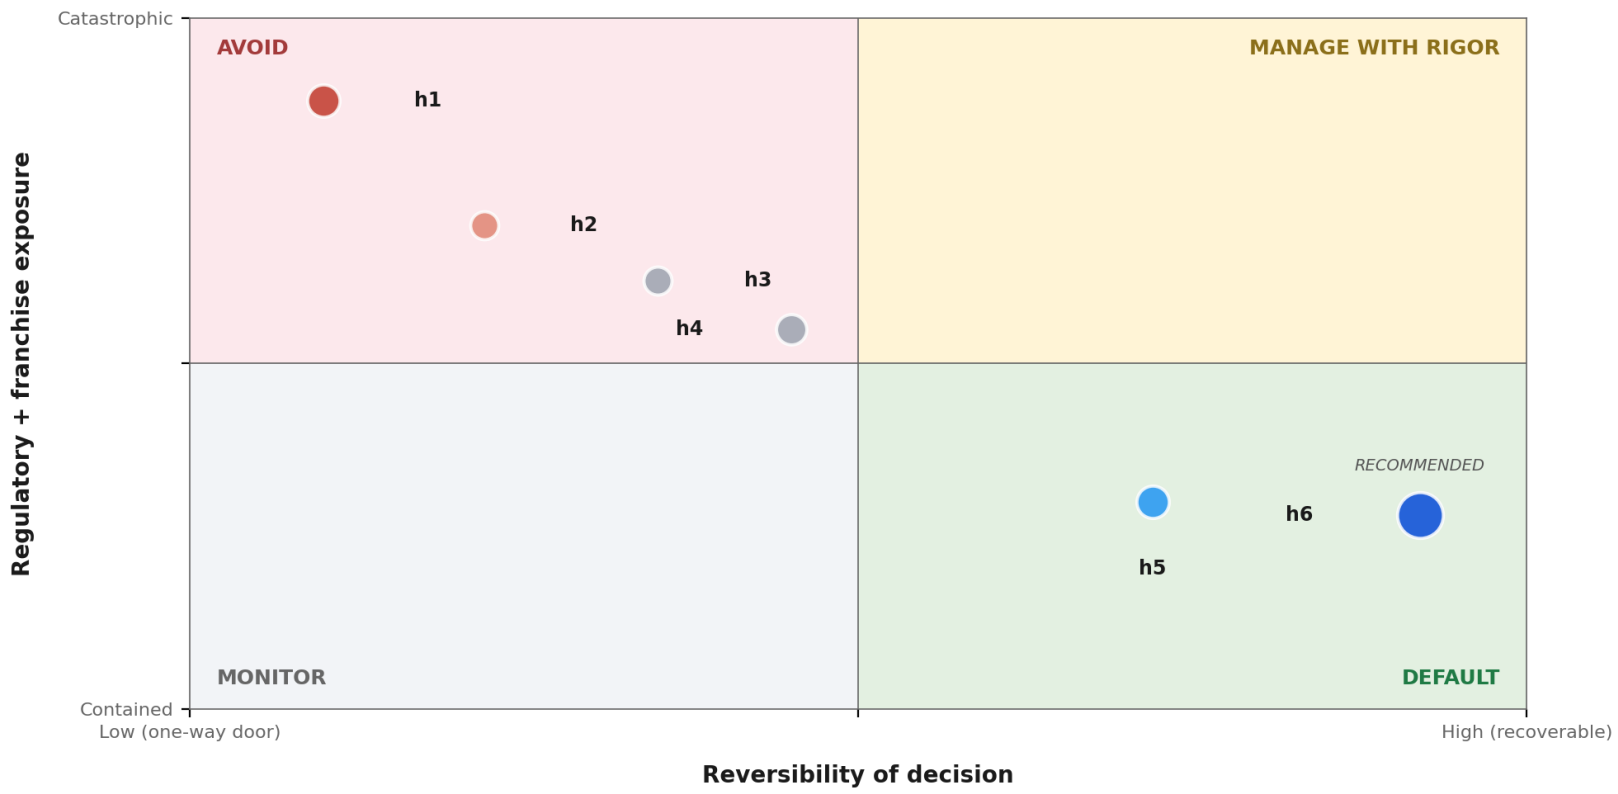

READ

- h1 — Universal optional — irreversible policy + Title IV/F-1 catastrophe
- h2 — Social product collapse — risk path of h1, same quadrant
- h3 — Faculty revolt — triggered by unilateral change
- h4 — Bifurcation — structural, mid-tier vulnerable
- h5 — Already optional — true for lectures, not labs/seminars
- h6 — Phased + structured pilot — preserves optionality, defuses 8/12 adversarial scenarios

Source: Synthesis from Decision Brief — bubble size = council_agreement × probability of dominance

PART 3

What the council found

Where the personas converged, where they diverged, and what the disagreement itself reveals about the decision.

The strongest signals point to regulatory catastrophe under universal-optional and structural antifragility under the phased pilot

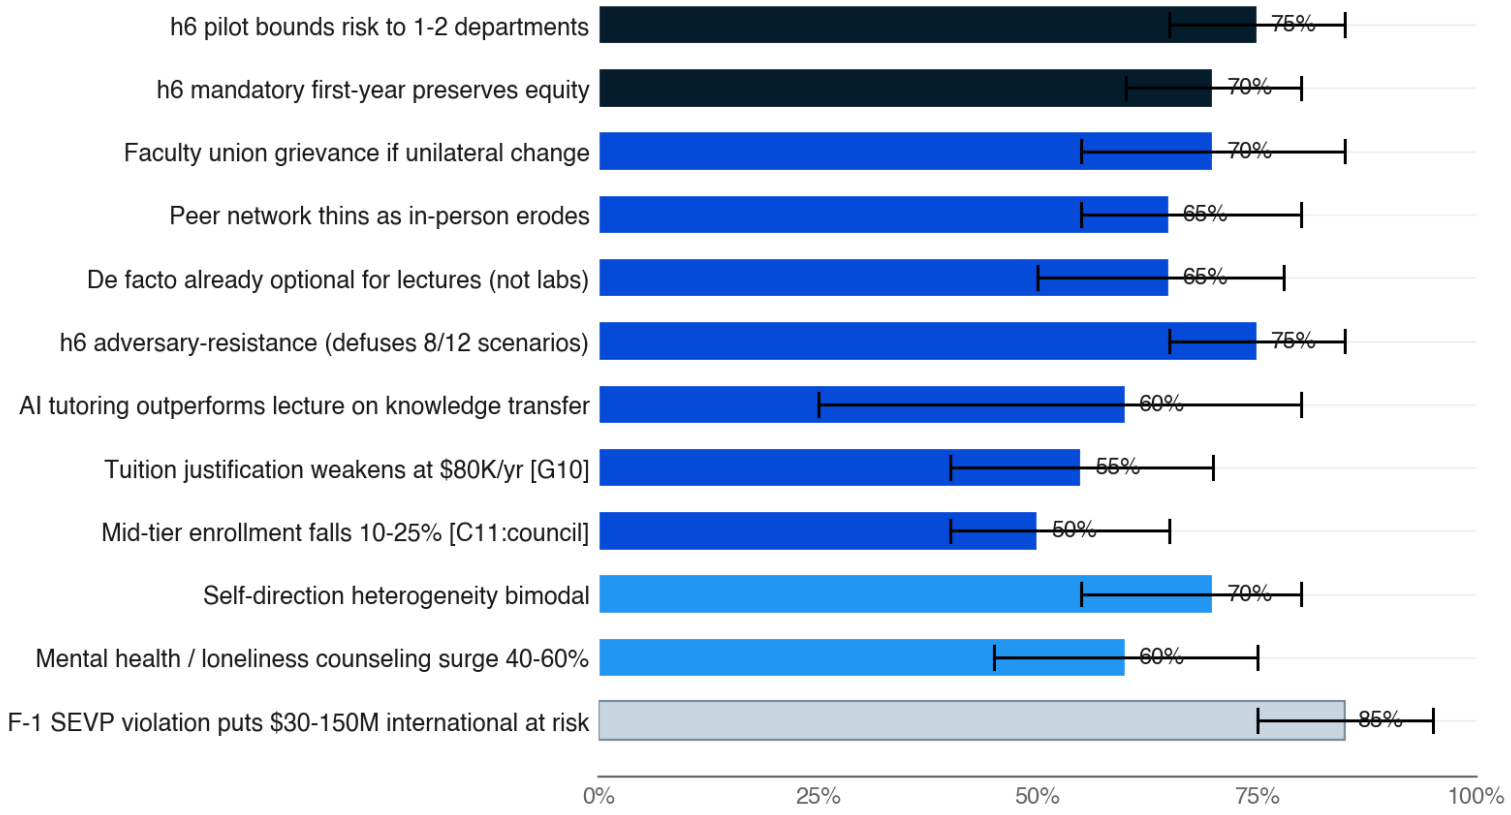

- Top 12 effects ranked by council_agreement × probability
- Bar color = council agreement (navy = 5/5, fading to pale = 1/5)
- Black bracket = persona disagreement range — width is the uncertainty signal
- Two highest-agreement effects (5/5) are h6 structural protections — the pilot bounds risk and mandatory first-year preserves equity
- Regulator's F-1 SEVP finding (1/5, P=0.85) — single-persona but existential; drove h6 dual-track compliance design

Source: Decision Brief — Effects Map (top 12 by council_agreement × probability) · [C3:regulator], [C9:council], [C13:iter2council]

The deepest disagreement — Harvard RCT replicability — IS the data

STRONGEST: REGULATOR (#1 PEER RANK)

- Won on specific regulatory mechanisms — F-1 SEVP one-online-course rule (P=0.85 [C3:regulator]); Title IV RSI reclassification ($50-300M/yr [C4:regulator] aid at risk); FERPA AI-vendor DPA; accreditor substantive-change trigger
- Each finding is potentially existential — universal optional is self-enforcing collapse without structured workaround
- Surfaced the dual-track compliance filing + required in-person seminar components that became h6's regulatory engine
- h6 protected by structuring: regulatory exposure shifts from catastrophic to bounded

DEEPEST DIVIDE: HARVARD RCT REPLICABILITY

- Optimist: 60% [C1:optimist] AI-hybrid replicates at institution scale (Harvard RCT [G1] effect size 0.7-1.3 SD)
- Pessimist + Customer: 25% [C1:pessimist] — N=194 lab study won't generalize across 30K students; self-direction heterogeneity bimodal
- Range 25-60% on h1 conditional outcome — disagreement IS the data, not noise to average away
- Optimist weakest (#5 peer rank) — overweighted Harvard sample, underweighted self-direction variance flagged by Customer
- BLIND SPOT caught: equity. Optional attendance is REGRESSIVE inside one tuition tier — high-achievers gain, first-gen lose, no price reduction

Source: Decision Brief — Council Dynamics · [C1:optimist], [C1:pessimist], [C2:customer]

Regulator scored highest on structural depth and counter-factual rigor; Optimist's narrow ranges drove its #5 peer rank

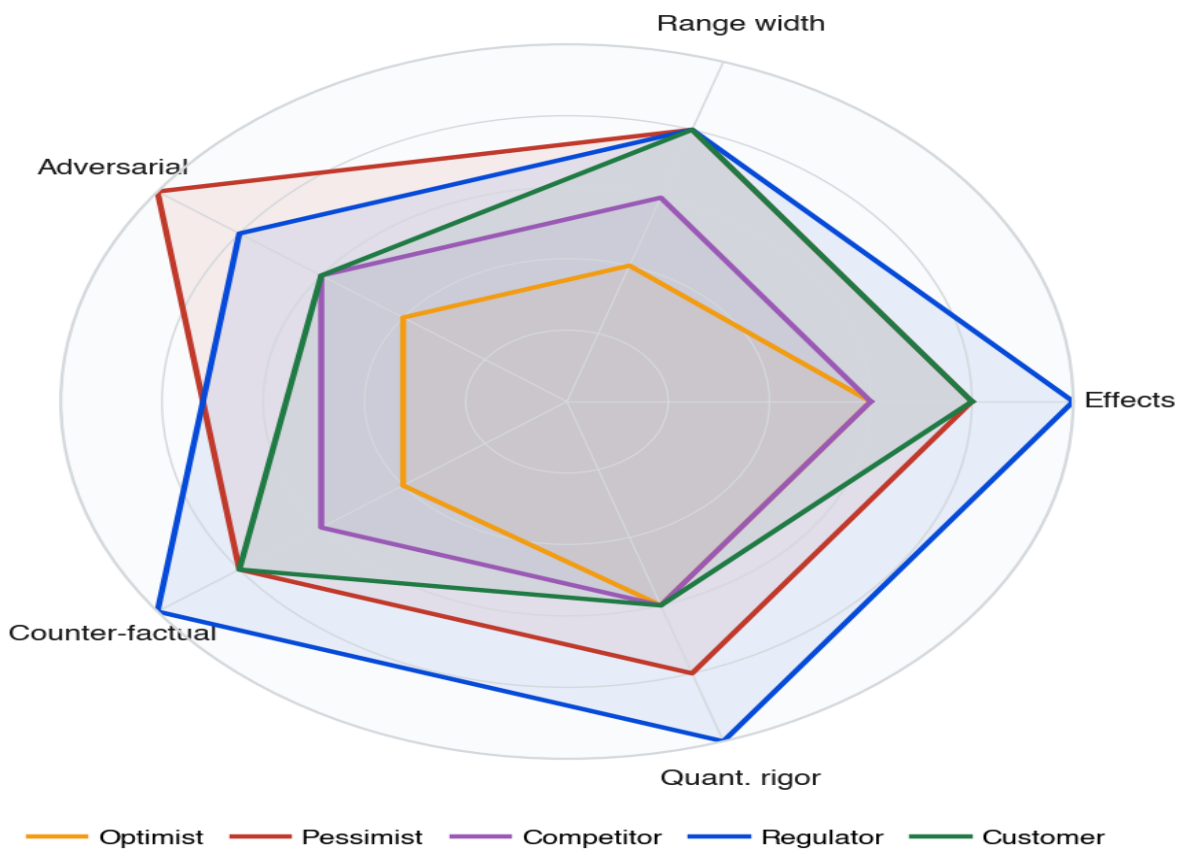

- Each persona scored 1-5 across five rigor dimensions, derived from peer-review and critique outputs
- Regulator's structural depth surfaced both regulatory mechanisms that synthesized into h6 — dual-track compliance and required in-person seminars
- Customer caught the equity blind spot — first-gen dropout amplification under universal optional
- Optimist's narrow ranges and weak counter-factual depth flagged by 4/4 reviewers — Harvard-RCT generalization optimism
- Pessimist + Regulator together cover the full adversarial surface; their alignment on regulatory irreversibility drove h6's safeguards

Source: Synthesized from Decision Brief — Council Dynamics + Critique outputs

Twelve adversarial scenarios — h6's structured design defuses 8 of them; 4 remain (foundation model, Ivy preempt, campus event, ED admin)

WORST CASES

- Title IV RSI reclassifies as correspondence; $50-300M/yr aid lost (P=0.40)
- F-1 SEVP mass violation; 15-25% international revenue drop (P=0.45)
- First-gen cohort 30%+ dropout year 1 — equity disaster + Title VI (P=0.35)
- Accreditor show-cause letter; focused review; reputational damage (P=0.25)

BLACK SWANS

- ChatGPT-6 dramatically better — lecture obsolete, accelerates transition (P=0.30)
- Ivy goes optional first; mid-tier first-mover narrative destroyed (P=0.08)
- Campus shooting / major event requires in-person community (P=0.10)
- ED admin change loosens Title IV; regulatory pressure removed (P=0.35)

IRRATIONAL ACTORS

- Student influencer (5M followers) attacks policy; applications drop (P=0.25)
- Sports alum boosters cut $10M [C34:council] donations (P=0.20)
- Faculty union strikes mid-semester; operations chaos (P=0.15)
- State legislature passes 'must require in-person' law (P=0.15)

Source: Decision Brief — Adversarial Scenarios · [C29:council] through [C36:council]

Four high-sensitivity assumptions are load-bearing — if any fails, the universal-optional path becomes indefensible

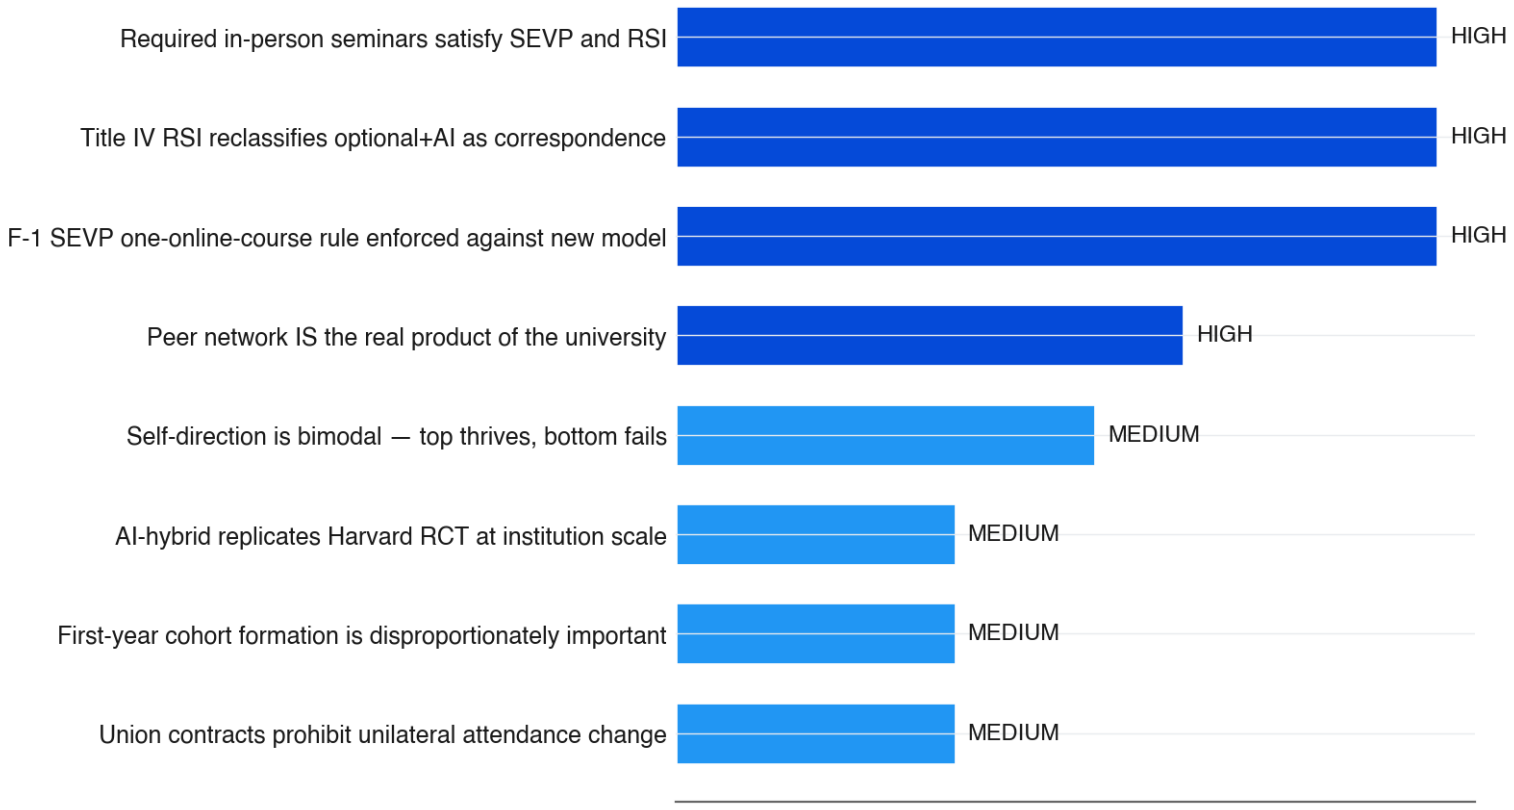

- Bar length = sensitivity weight × number of effects impacted
- Dark-blue bar = HIGH sensitivity, lighter = MEDIUM
- Three of four HIGH-sensitivity assumptions are external regulatory interpretations — the university controls none directly
- h6's dual-track filing converts 'F-1 SEVP enforced' from catastrophic to bounded: the structure pre-files in-person components that satisfy the rule
- If 'required in-person seminars satisfy SEVP+RSI' interpretation fails on legal review (P=0.30), h6's structural workaround fails — pilot must be paused, not the policy reversed

Source: Decision Brief — Key Assumptions · sensitivity per effects-impacted × HIGH/MEDIUM/LOW weight

PART 4

Path forward

Seven gates with kill criteria — every commitment is bounded by faculty, regulator, and pilot data.

The recommendation maps to a seven-gate roadmap with explicit kill criteria at every stage

Every gate has a binding kill criterion; failure at any gate triggers either pause, redesign, or full rollback

When	Action	Decision point	Kill criteria
Q1 2026	Faculty senate briefing + union engagement	Senate accepts pilot design	Senate rejects → pause
Q2 2026	Substantive change filing + dual-track compliance design	Accreditor approves path	Accreditor blocks → redesign
Q3 2026	SEVP pre-filing + DOE consultation on RSI	Structure clears regulatory review	Regulators block → abandon pilot
Q4 2026	Pilot launch in 1-2 departments; mandatory first-year intact	Opt-in pilot enrolls 300-500 students	Enrollment <150 → insufficient power
2027-28	Annual outcome reporting: outcomes, retention, equity	Year-2 data shows positive indicators	Major negative trend → halt
Q4 2028	3-year pilot review with external advisory board	Outcomes support institution rollout	Pilot fails → end pilot, maintain current
2029+	Phased institution-wide rollout if pilot succeeds	Board + senate + union approval	First-gen retention drops >5% → rollback

Source: Decision Brief — Appendix A: Decision Timeline · phased gates per [C13:iter2council], [C14:iter2council]

Selected sources

Tags appear inline throughout the deck — [G#] = ground data, [C#:persona] = council finding, [C#:adversary] = adversarial scenario

Tag	Type	Claim	Source
G1	Ground	Harvard 2025 RCT: AI tutoring effect size 0.73-1.3 SD vs active learning	nature.com
G3	Ground	Human-AI hybrid > AI-only > classroom-only	phys.org
G6	Ground	Post-COVID absence rates 73-83% in some courses	timeshighereducation.com
G10	Ground	Tuition: public $10K/$26K, private $37K, elite $80K total cost	research.com
G15	Ground	43% students used ChatGPT; 89% for homework; 8/10 used GenAI 2025	nerdynav.com
G21	Ground	Minerva: 100% online + active learning + rotating cohort; elite alt	general
C1:optimist	Council	AI-hybrid replicates at scale P=0.60 (vs Pessimist 0.25)	iter-1 optimist
C2:customer	Council	Self-direction heterogeneity bimodal P=0.70 — bimodal failure mode	iter-1 customer
C3:regulator	Council	F-1 SEVP one-online-course violation P=0.85; $30-150M international	iter-1 regulator
C4:regulator	Council	Title IV RSI reclassification; $50-300M/yr federal aid at risk	iter-1 regulator
C13:iter2council	Council	H6 pilot bounds risk P=0.75 (5/5 council)	iter-2 council
C14:iter2council	Council	H6 mandatory first-year preserves equity P=0.70 (5/5 council)	iter-2 council
C15:iter2council	Council	H6 adversary-resistance — defuses 8/12 adversarial scenarios	iter-2 council

Source: Decision Brief — Sources (full list in companion brief markdown)