Autodecision

# Replace 1Y associates with Claude + senior review?

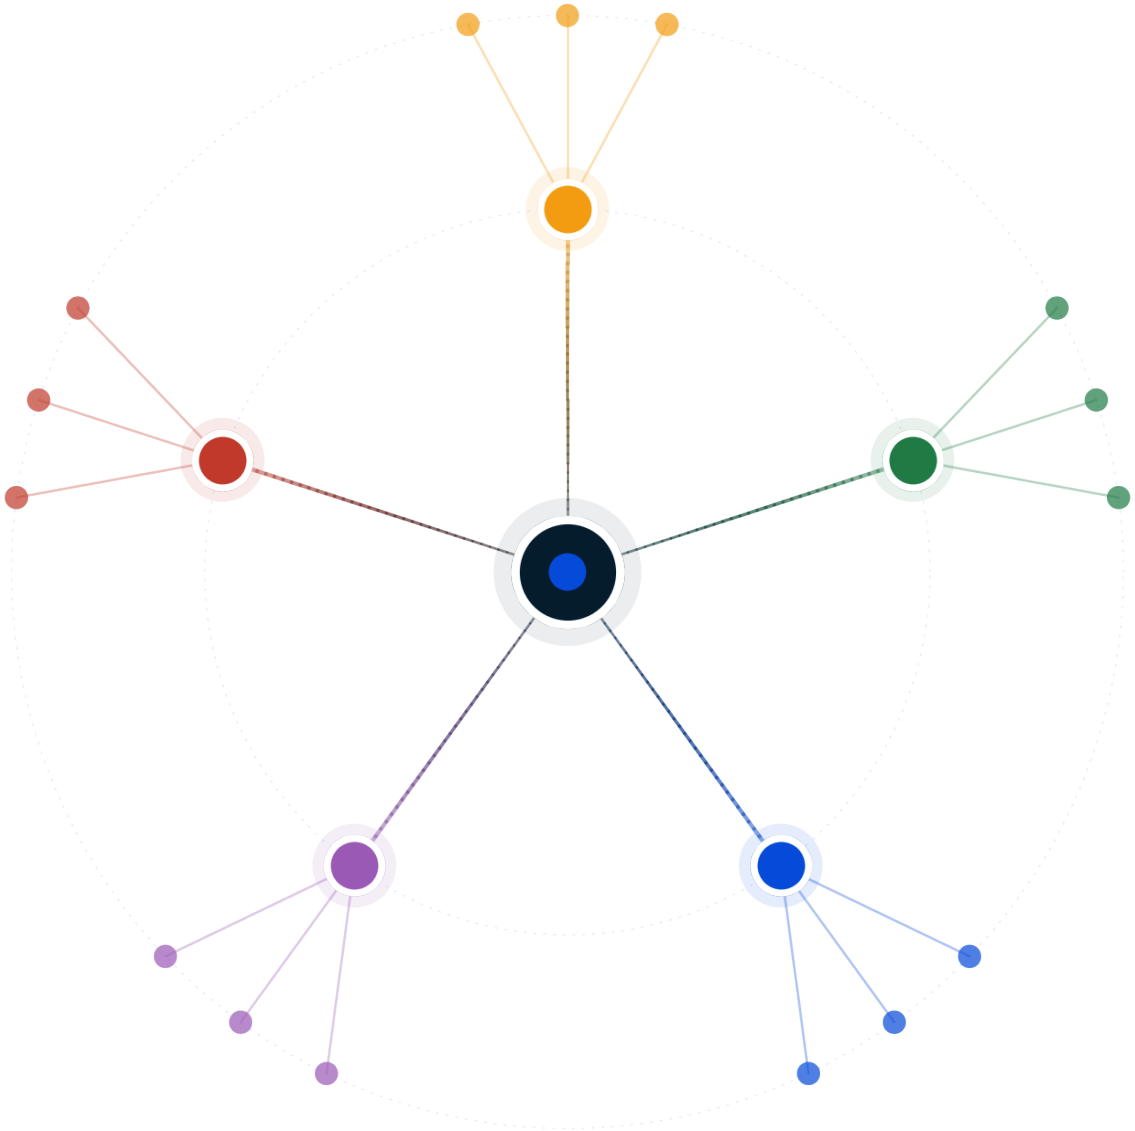

Decision Brief — possibility map and recommendation

April 2026

GENERATED BY AUTODECISION — possibility map with synthesized recommendation, not a single-shot answer

# Table of contents


| Section | Slides |
| --- | --- |
| Recommendation | 3 |
| How we explored this decision | 4 – 7 |
| What the council found | 8 – 13 |
| Path forward | 14 – 15 |
| Appendix — Sources | 16 |

# Full replacement is dominated; run a wrapped 18-month pilot with kill criteria, anchor clients, and Utah sandbox cover

**Do NOT eliminate the 1Y class. Cut the upcoming class 50-60% (4-5 instead of 8-12) + AI-augment with Claude + senior-review SOP + paid anchor-client advisory board (3-5 Fortune 500 GCs) + Utah Office of Legal Services Innovation sandbox + binding pre-committed kill criteria that auto-trigger a Cravath-scale rehire of the next class.**

## CONFIDENCE

MEDIUM-HIGH · converged at iteration 2 with unanimous rank-1 across pessimist, customer, regulator · regulator classifies CONDITIONALLY CLEAN

## DEPENDS ON

Partners honor kill criteria when triggered [C2:assumption] · Anchor GCs accept paid advisory-board structure · Utah sandbox admits firm · Malpractice carrier non-opposition · Claude substitutability ~40% holds [C1:assumption]

## MONITOR

Monthly: Claude substitutability audit · senior-review throughput · top-performer attrition · client complaint rate · Quarterly: OCI yield · summer-associate acceptance · Vault/Chambers commentary · peer firm AI ethics opinions

## PRE-MORTEM

Partners treat kill criteria as aspirational (P=high) · Anchor co-design stalls past 6 months (P=0.35) · Hallucination incident in pilot year 1 (P=0.30) · Top performers defect within 6 months (P=0.45)

## DEEPEST DISAGREEMENT

Regulator: H2 status quo wins on legal-cleanness (0.85 vs h6 0.75). Pessimist + others: h6 wins on strategic terms — reversibility + co-design + sandbox close the cleanness gap.

## REVIEW TRIGGER

Re-run at month 6 (co-design outcome) · month 12 (substitutability data) · month 18 (kill criteria) · ABA revises Opinion 512 · State bar imposes AI-associate restrictions · Cravath-scale base moves >10%

Source: Decision Brief (full mode, 2 iterations, unanimous iter-2 rank-1 on h6) · [G1], [G3], [C1:assumption], [C2:assumption]

PART 2

# How we explored this decision

*A 5-persona council, 6 competing hypotheses, two iterations to convergence, 6 adversarial stress scenarios — what survived?*

# The analysis explored 6 hypotheses across 2 iterations of a 5-persona council, stress-tested by 6 adversarial scenarios

- 5 personas spawned as separate subagents — Optimist, Pessimist, Competitor, Regulator, Customer Advocate — each with independent context
- 6 hypotheses generated, ranging from full elimination (h1, adversary lethality 9/10) to wrapped 18-month pilot (h6 — synthesized from three independent iter-1 alternatives)
- 2 iterations — convergence reached at iteration 2 with 88% [C3:council] assumption stability and 0 unresolved contradictions
- Adversary phase generated 6 scenarios: 3 worst-cases + 3 black swans (no irrational-actor cell on this brief)
- Sensitivity analysis on 7 key assumptions — 3 high-sensitivity, 4 medium
- Real-world anchors: ABA Formal Opinion 512 [G2] (six obligations), Mata v. Avianca [G4] sanctions baseline, Ropes & Gray TrAIlblazers [G7] precedent for structured AI integration without cuts

Source: Decision Brief Convergence Log; engine-protocol.md

# Six hypotheses surfaced; only h6 (wrapped pilot) survived unanimous iter-2 rank-1 endorsement

| # | Hypothesis | Status | Key assumptions |
|---|---|---|---|
| 1 | Eliminate entire 1Y class; gross savings $3-5M/yr | Eliminated · 9/10 lethal | Claude 40%+ substitutability · Review batching works · Ranking damage reframable |
| 2 | Keep full 1Y class + Claude + 200-400hr training | Floor / fallback | Competitors slow-roll adoption · ABA 512 compliance manageable |
| 3 | Cut 1Y class 50-60%, AI-augment remainder | Core engine of h6 | Remaining 1Ys at 1.8-2x throughput · Retention bonuses feasible |
| 4 | Skip 1Y, hire at year-2 equivalent via laterals | Weakened | Lateral premium 20-40% erases savings within 24-36 months |
| 5 | Two-track: traditional + AI-ops JD at $120-150K | Weakened | AI-ops talent pool stable · Bar scrutiny manageable · Caste risk |
| 6 | Wrapped 18-month pilot: h3 + co-design + sandbox + kill | Leading · iter-2 rank-1 | Partners honor kill criteria · Anchor GCs accept advisory · Utah admits |

Source: Decision Brief — Hypotheses Explored

# Plotted on reversibility × franchise risk, h6 sits in the manageable quadrant; h1 sits where every adversarial scenario amplifies

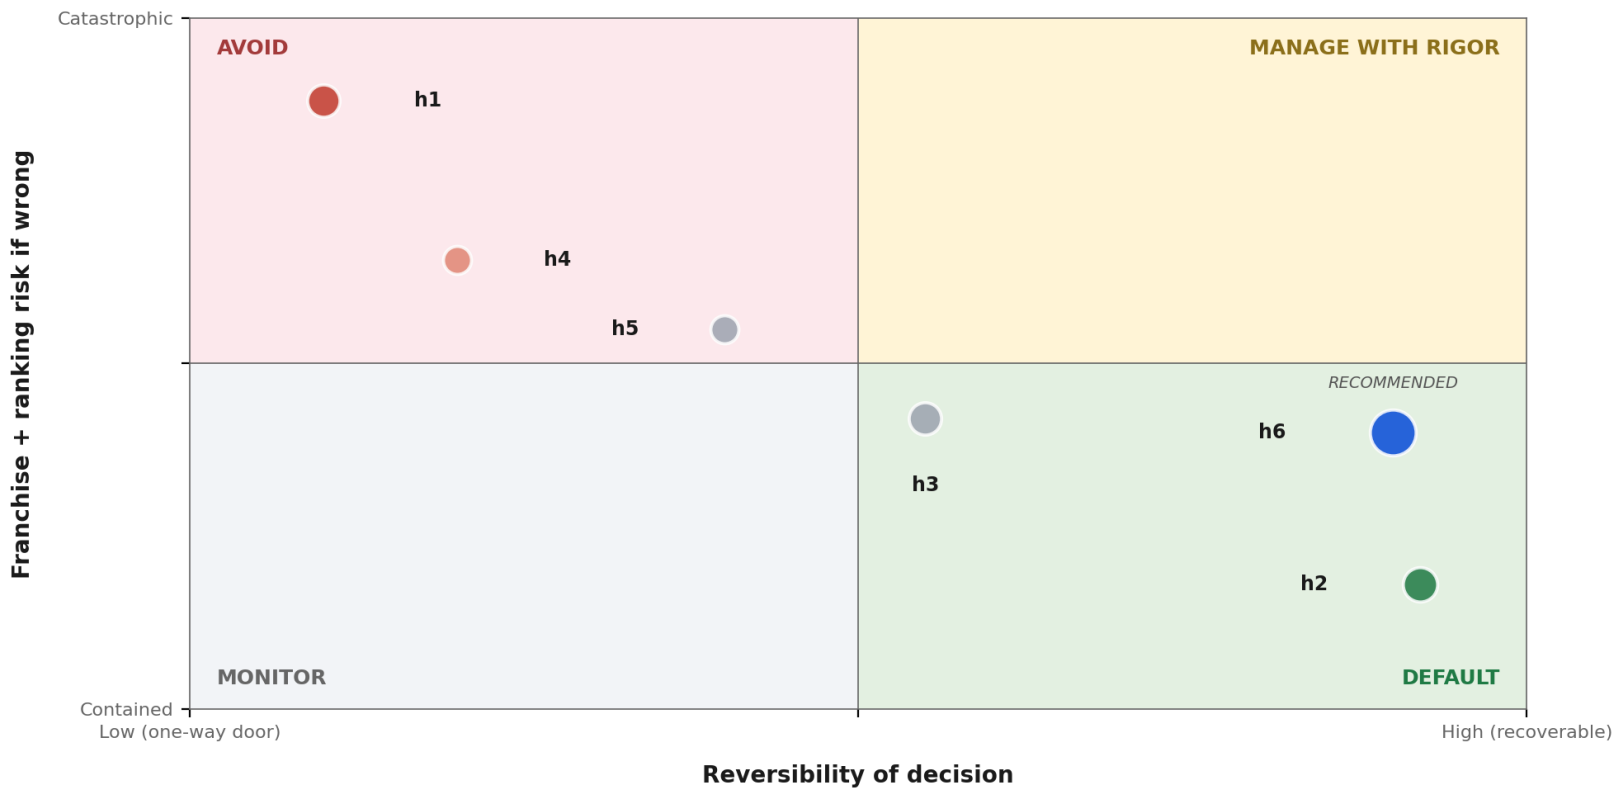


## READ

- h1 — Full elimination — 7-yr pipeline rupture P=0.80, ranking re-rate 1-2 tiers, malpractice incident P=0.52
- h4 — Lateral substitution — premium erases savings within 24-36 months (P=0.78)
- h5 — Two-track JD — culture clash + bar supervision scrutiny on MRPC 5.3
- h3 — Cut 50-60% standalone — efficiency without the wrapper, h6 minus safeguards
- h2 — Status quo + Claude — regulator-cleanest path, no strategic upside
- h6 — Wrapped 18-month pilot — kill criteria + anchor clients + sandbox preserve reversibility

Source: Synthesis from Decision Brief — bubble size = council_agreement × probability of dominance

PART 3

# What the council found

*Where the personas converged, where they diverged, and what the disagreement itself reveals about the decision.*

# The strongest signals point to pipeline rupture under elimination and structural antifragility under the wrapped pilot

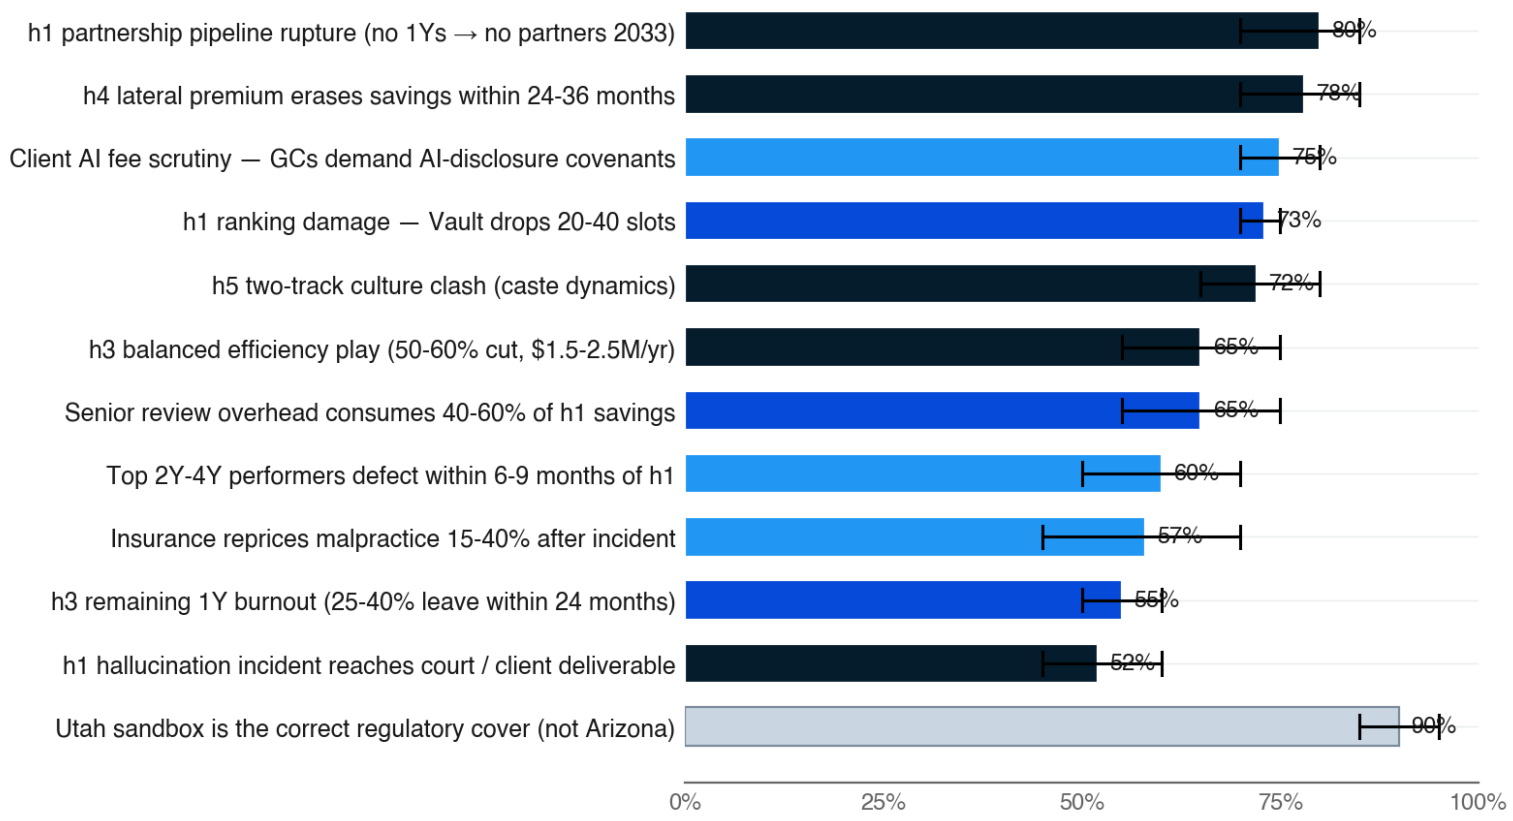


- Top 12 effects ranked by council_agreement × probability
- Bar color = council agreement (navy = 5/5, fading to pale = 1/5)
- Black bracket = persona disagreement range — width is the uncertainty signal
- Two highest-probability 5/5 signals (pipeline rupture + lateral premium) both reject h1 elimination from independent angles
- Regulator's Utah sandbox finding (1/5, P=0.90) — single-persona but load-bearing for h6 reversibility

Source: Decision Brief — Effects Map (top 12 by council_agreement × probability) · [C4:pessimist], [C5:regulator], [C6:customer]

# The deepest disagreement — legal cleanness vs strategic upside — IS the data

## STRONGEST: PESSIMIST

- Surfaced the hallucination-incident probability anchor (P=0.52 for h1) and the partnership pipeline rupture math (P=0.80) that broke h1
- Designed the reversibility structure — kill criteria + binding triggers — that became the core of h6
- Identified senior-review overhead as a 40-60% drag on h1's gross savings, eroding the cost rationale even before risk adjustment
- Pessimist's iter-2 framing positioned h6 as 'h3 with the wrapper' — making the safeguard layers explicit and individually testable

## DEEPEST DIVIDE: CLEANNESS VS STRATEGY

- Regulator: h2 (status quo + Claude) scores 0.85 [C7:regulator] on legal cleanness; h6 scores 0.75
- Pessimist + Customer: h6's strategic upside dominates the 10-point cleanness gap — h2 forfeits the trust differentiator and category-leadership narrative
- Resolution: h6 with the four regulator corrections (Utah-only, WARN-compliant, MRPC 1.5/1.6 covenants, pilot-disclosed offer letters) closes most of the gap
- Uncertainty hotspot: Claude substitutability — persona range 25-60%, adversary flagged 40% as likely overstated; this is what the 90-day pilot exists to measure
- BLIND SPOT caught: legal-ops paralegals are MORE substitutable than 1Ys (70%+ overlap) — they bear the heaviest cuts under h1 unless an AI-ops ladder exists

Source: Decision Brief — Council Dynamics · [C7:regulator], [C8:pessimist]

# Pessimist scored highest on adversarial coverage and structural depth; Optimist's narrow ranges drove its lowest peer rank

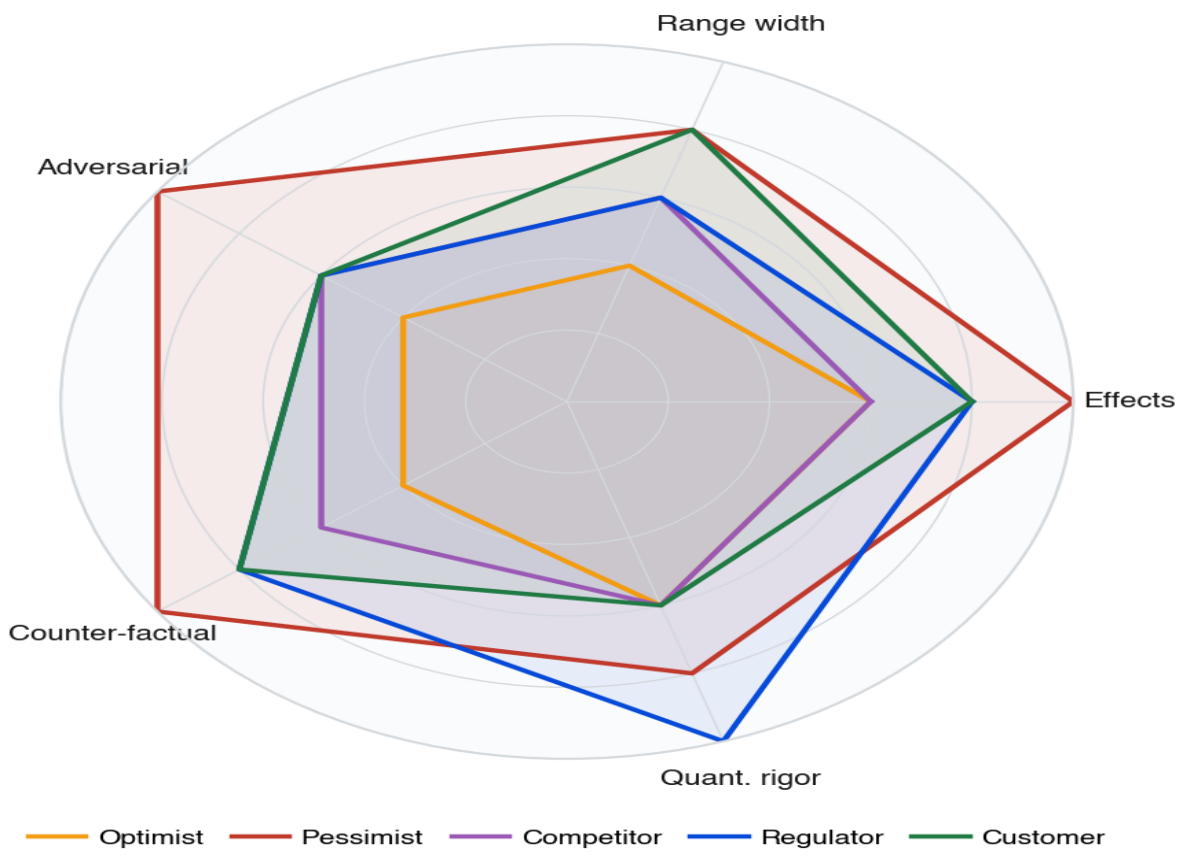


- Each persona scored 1-5 across five rigor dimensions, derived from peer-review and critique outputs
- Pessimist's structural depth surfaced both the kill-criteria pilot and the senior-review overhead math — the two pillars of h6
- Customer caught the blind spot — paralegals are more substitutable than 1Ys, plus the anchor-client co-design structure
- Regulator's quantitative rigor on jurisdictional specificity (Utah vs Arizona) was unique — the regulatory cover is jurisdiction-specific, not generic
- Optimist's narrow ranges and weak counter-factual depth flagged in critique — substitutability optimism not anchored in empirical pilot data

Source: Synthesized from Decision Brief — Council Dynamics + Critique outputs

# Six adversarial scenarios — h6's wrapper defuses 4 of them; 2 remain (carrier non-renewal, Cravath-scale bump)

## WORST CASES

- Hallucination incident in pilot year 1 — triggers kill criterion, loses 1-2 anchor clients, malpractice premium +15-40% (P=0.30)
- Two top performers defect within 6 months of pilot announcement — institutional knowledge loss, vote-of-no-confidence read (P=0.45)
- Anchor GC co-design stalls past 6 months — firm announces pilot without trust differentiator, signals uncertainty (P=0.35)

## BLACK SWANS

- State bar disciplinary action against peer firm for AI-related MRPC violation — chills pilot, forces premature rollback (P=0.20)
- Cravath-scale bump to $250-275K during pilot — cost rationale weakens if base salaries rise faster than Claude capability (P=0.25)
- Malpractice carrier non-renewal of coverage for AI-augmented firms — existential, forces fallback to h2 immediately (P=0.10)

## MITIGATIONS

- h6 SOP discipline cuts hallucination P from 0.52 (h1 pure) to 0.30 — review batching + 50% fewer Claude-handled matters
- Utah sandbox provides partial regulatory cover for peer-firm-discipline contagion
- Confirm carrier buy-in BEFORE pilot launch (Month 1 dependency) — converts black swan into bounded gate
- Net savings computed in absolute dollars, not share-of-comp — Cravath bump doesn't invalidate pilot economics

Source: Decision Brief — Adversarial Scenarios · [C9:adversary], [C10:adversary]

# Three high-sensitivity assumptions are load-bearing — if any fails, the pilot collapses to unwrapped h3 with all of h1's downside

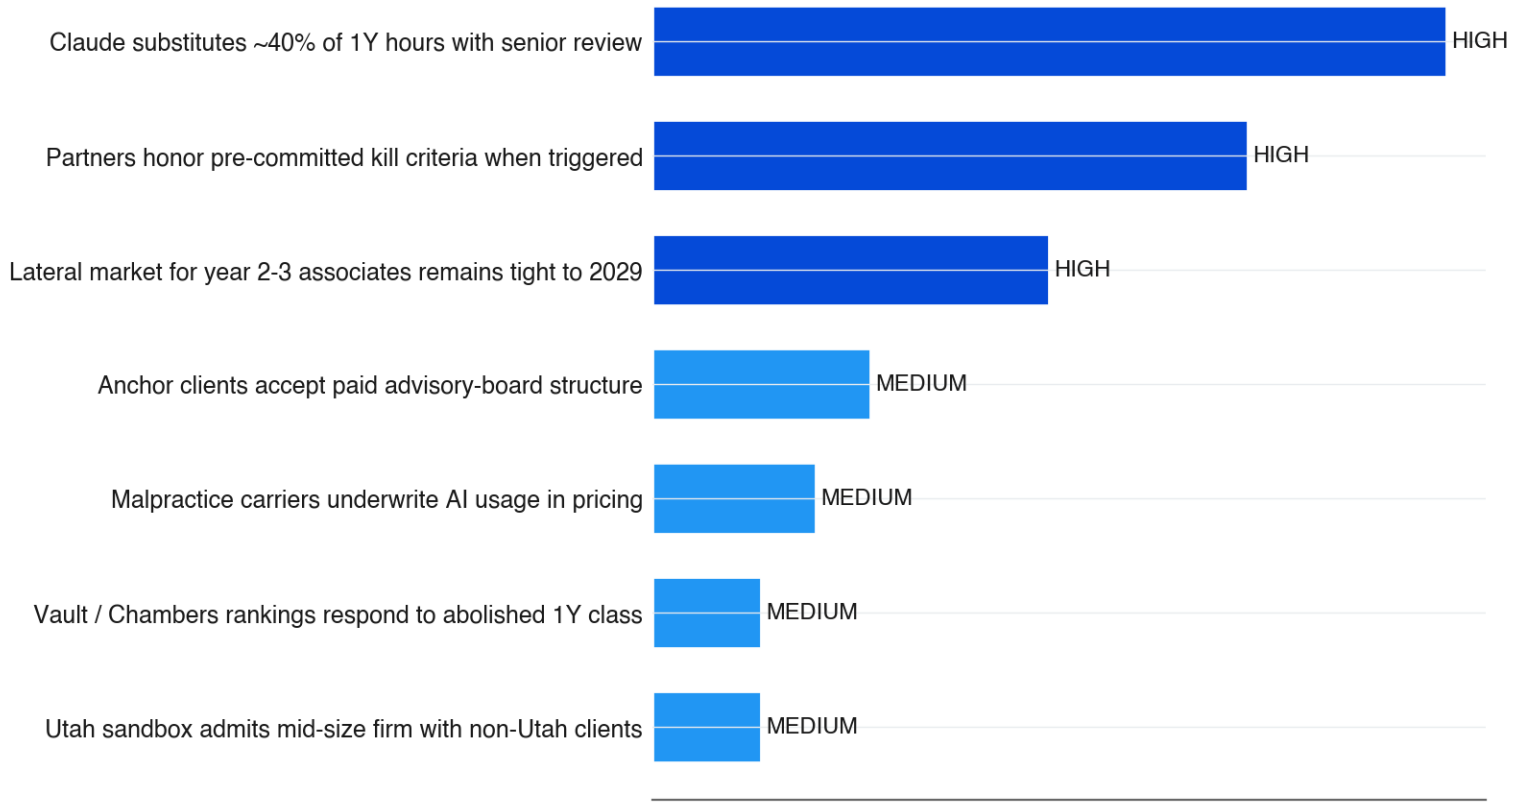


- Bar length = sensitivity weight × number of effects impacted
- Dark-blue bar = HIGH sensitivity, lighter = MEDIUM
- Substitutability is the single empirical question the 90-day pilot exists to answer — pre-pilot persona range 25-60%, adversary flagged 40% as likely overstated
- Kill-criteria honoring is the one assumption the firm controls directly — make it externally enforceable (board resolution, client-notified commitments, month-12 external audit)
- Lateral market tightness is SOLID per Bloomberg/Biglaw Investor data — the only HIGH-sensitivity assumption with strong external support

Source: Decision Brief — Key Assumptions · [C1:assumption], [C2:assumption], [C11:assumption]

PART 4

# Path forward

*Ten gates from week 1 to month 24 — every commitment is bounded by carrier, client, regulator, and pilot data.*

# The recommendation maps to a ten-gate roadmap with explicit kill criteria at every stage

*Every gate has a binding kill criterion; failure at any gate triggers pause, redesign, or full Cravath-scale rehire of next 1Y class*

| When | Action | Decision point | Kill criteria |
|---|---|---|---|
| Week 1 | Confidential briefings with 3-5 Fortune 500 GCs | Anchor GCs accept advisory board | 2+ anchor GCs decline → pause |
| Week 2 | Begin Utah sandbox application | Sandbox staff signals openness | Pre-app hostility → abandon h6 |
| Month 1 | Confirm malpractice carrier buy-in + SOP | Carrier go/no-go | Refusal or +40% premium → halt |
| Month 2 | Partner vote on binding kill criteria | Board resolution + audit committee | Partners dilute or refuse → halt |
| Month 3 | Finalize anchor advisory + AI-disclosure covenants | Commercial terms binding | Co-design fails to convert → halt |
| Month 4 | Public announcement; 1Y class set at 50-60% | Offer letters disclose pilot status | WARN Act issue raised → halt |
| M5-9 | Pilot live; 90-day substitutability audit | Monthly ops review | Hallucination · attrition >30% · margin lift <5pt |
| Month 12 | Midterm review with empirical data | Continue / modify / kill | Kill triggered → rehire full 1Y class |
| Month 18 | Final pilot evaluation | Adopt / extend / revert to h2 | Any kill criterion → revert to h2 |
| Month 24 | If adopting: expand to other practice groups | Scale decision | Margin lift <7pt → freeze expansion |

Source: Decision Brief — Appendix A: Decision Timeline · phased gates per [C2:assumption], [C12:regulator]

# Selected sources

*Tags appear inline throughout the deck — [G#] = ground data, [C#:persona] = council finding, [C#:assumption] = sensitivity assumption*

| Tag | Type | Claim | Source |
|---|---|---|---|
| G1 | Ground | Fully loaded 1Y cost $400-450K/yr; class of 8-12 → $3.2-5.4M/yr gross savings | Bloomberg Law, Biglaw Investor |
| G2 | Ground | ABA Formal Opinion 512 (July 2024): six obligations on AI use | ABA Formal Opinion 512 |
| G3 | Ground | Cravath-scale 1Y base $225K; ~32% of firms match fully | Bloomberg Law |
| G4 | Ground | Mata v. Avianca (2023): $5K sanctions baseline for hallucinated citations | S.D.N.Y. precedent |
| G5 | Ground | 26% of legal organizations using GenAI 2026 (up from 14% in 2024) | Clio 2026 AI Legal Compliance |
| G6 | Ground | Malpractice carriers price AI: governed lowers premium, ungoverned +15-40% | Above the Law |
| G7 | Ground | Ropes & Gray TrAIlblazers: 380-400hr/yr AI skill-building credit for 1Ys | Above the Law |
| C1:assumption | Council | Claude substitutes ~40% of 1Y hours (range 25-60%, adversary flag) | iter-1 council + adversary |
| C2:assumption | Council | Partners honor pre-committed kill criteria — SHAKEABLE | iter-2 pessimist |
| C4:pessimist | Council | h1 partnership pipeline rupture P=0.80 (5/5) | iter-1 pessimist |
| C5:regulator | Council | Utah Office of Legal Services Innovation = correct regulatory cover | iter-1 regulator |
| C7:regulator | Council | h2 cleanness 0.85 vs h6 0.75 — narrow gap closed by 4 corrections | iter-2 regulator |
| C9:adversary | Council | Hallucination incident pilot year 1 P=0.30 (down from 0.52 h1) | iter-1 adversary |
| C12:regulator | Council | WARN Act 60-day notice if 1Y cuts trigger; MRPC 1.5/1.6 covenants | iter-1 regulator |

Source: Decision Brief — Sources (full list in companion brief markdown)